

exploration

August 19, 2026

1 C-MAPS FD001 — Exploratory Data Analysis

Sensor variance analysis and degradation pattern exploration for the FD001 turbofan dataset.

```
[23]: import sys
      sys.path.append('../src')

      import pandas as pd
      import numpy as np
      import matplotlib.pyplot as plt
      import seaborn as sns

      from data_loader import load_data

      sns.set_style('whitegrid')
      plt.rcParams['figure.figsize'] = (12, 6)

      %matplotlib inline
```

1.1 1. Load Data

```
[24]: train_df, test_df, test_rul = load_data('../data/')

      print(f"Training set: {train_df.shape}")
      print(f"Test set: {test_df.shape}")
      print(f"True RUL values: {test_rul.shape}")
```

```
loaded training data: 20631 records from 100 engines
calculated RUL: min=0, max=125, mean=86.8
loaded test data: 13096 records from 100 engines
loaded ground truth RUL: min=7, max=125, mean=74.5
Training set: (20631, 27)
Test set: (13096, 26)
True RUL values: (100, 1)
```

1.2 2. Data Overview

```
[25]: train_df.head()
```

```
[25]:
```

	unit_id	cycle	op_setting_1	op_setting_2	op_setting_3	sensor_1	\
0	1	1	-0.0007	-0.0004	100.0	518.67	
1	1	2	0.0019	-0.0003	100.0	518.67	
2	1	3	-0.0043	0.0003	100.0	518.67	
3	1	4	0.0007	0.0000	100.0	518.67	
4	1	5	-0.0019	-0.0002	100.0	518.67	

	sensor_2	sensor_3	sensor_4	sensor_5	...	sensor_13	sensor_14	\
0	641.82	1589.70	1400.60	14.62	...	2388.02	8138.62	
1	642.15	1591.82	1403.14	14.62	...	2388.07	8131.49	
2	642.35	1587.99	1404.20	14.62	...	2388.03	8133.23	
3	642.35	1582.79	1401.87	14.62	...	2388.08	8133.83	
4	642.37	1582.85	1406.22	14.62	...	2388.04	8133.80	

	sensor_15	sensor_16	sensor_17	sensor_18	sensor_19	sensor_20	\
0	8.4195	0.03	392	2388	100.0	39.06	
1	8.4318	0.03	392	2388	100.0	39.00	
2	8.4178	0.03	390	2388	100.0	38.95	
3	8.3682	0.03	392	2388	100.0	38.88	
4	8.4294	0.03	393	2388	100.0	38.90	

	sensor_21	RUL
0	23.4190	125
1	23.4236	125
2	23.3442	125
3	23.3739	125
4	23.4044	125

[5 rows x 27 columns]

```
[26]: train_df.info()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 20631 entries, 0 to 20630
Data columns (total 27 columns):
#   Column          Non-Null Count  Dtype
---  -
0   unit_id         20631 non-null  int64
1   cycle           20631 non-null  int64
2   op_setting_1    20631 non-null  float64
3   op_setting_2    20631 non-null  float64
4   op_setting_3    20631 non-null  float64
5   sensor_1        20631 non-null  float64
6   sensor_2        20631 non-null  float64
```

```

7  sensor_3      20631 non-null  float64
8  sensor_4      20631 non-null  float64
9  sensor_5      20631 non-null  float64
10 sensor_6      20631 non-null  float64
11 sensor_7      20631 non-null  float64
12 sensor_8      20631 non-null  float64
13 sensor_9      20631 non-null  float64
14 sensor_10     20631 non-null  float64
15 sensor_11     20631 non-null  float64
16 sensor_12     20631 non-null  float64
17 sensor_13     20631 non-null  float64
18 sensor_14     20631 non-null  float64
19 sensor_15     20631 non-null  float64
20 sensor_16     20631 non-null  float64
21 sensor_17     20631 non-null  int64
22 sensor_18     20631 non-null  int64
23 sensor_19     20631 non-null  float64
24 sensor_20     20631 non-null  float64
25 sensor_21     20631 non-null  float64
26 RUL           20631 non-null  int64
dtypes: float64(22), int64(5)
memory usage: 4.2 MB

```

1.3 3. RUL Distribution

```


[27]: plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.hist(train_df['RUL'], bins=50, edgecolor='black', alpha=0.7)
plt.xlabel('Remaining Useful Life (cycles)', fontsize=12)
plt.ylabel('Frequency', fontsize=12)
plt.title('RUL Distribution (Training Set)', fontsize=14, fontweight='bold')
plt.grid(True, alpha=0.3)

plt.subplot(1, 2, 2)
plt.hist(test_rul['RUL'], bins=30, edgecolor='black', alpha=0.7, color='orange')
plt.xlabel('Remaining Useful Life (cycles)', fontsize=12)
plt.ylabel('Frequency', fontsize=12)
plt.title('RUL Distribution (Test Set)', fontsize=14, fontweight='bold')
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```




1.4 4. Engine Lifecycle Analysis

```
[28]: # Lifecycle length distribution
lifecycle_lengths = train_df.groupby('unit_id')['cycle'].max()

plt.figure(figsize=(10, 5))
plt.hist(lifecycle_lengths, bins=30, edgecolor='black', alpha=0.7,
        color='steelblue')
plt.xlabel('Engine Lifecycle Length (cycles)', fontsize=12)
plt.ylabel('Number of Engines', fontsize=12)
plt.title('Distribution of Engine Lifecycles', fontsize=14, fontweight='bold')
plt.grid(True, alpha=0.3)
plt.show()

print(f"Average lifecycle: {lifecycle_lengths.mean():.1f} cycles")
print(f"Min lifecycle: {lifecycle_lengths.min()} cycles")
print(f"Max lifecycle: {lifecycle_lengths.max()} cycles")
```



Average lifecycle: 206.3 cycles

Min lifecycle: 128 cycles


Max lifecycle: 362 cycles

1.5 5. Sample Engine Degradation

```
[29]: # Plot RUL degradation for first 5 engines
plt.figure(figsize=(14, 6))

for engine_id in range(1, 6):
    engine_data = train_df[train_df['unit_id'] == engine_id]
    plt.plot(engine_data['cycle'], engine_data['RUL'], label=f'Engine_{engine_id}', linewidth=2)

plt.xlabel('Cycle', fontsize=12)
plt.ylabel('Remaining Useful Life (cycles)', fontsize=12)
plt.title('RUL Degradation Curves - Sample Engines', fontsize=14, fontweight='bold')
plt.legend()
plt.grid(True, alpha=0.3)
plt.show()
```




1.6 6. Sensor Correlation Analysis

```
[30]: # Calculate sensor variances
sensor_cols = [col for col in train_df.columns if col.startswith('sensor_')]
sensor_variances = train_df[sensor_cols].var().sort_values(ascending=False)

plt.figure(figsize=(12, 6))
plt.bar(range(len(sensor_variances)), sensor_variances.values,
        edgecolor='black', alpha=0.7)
plt.xticks(range(len(sensor_variances)), sensor_variances.index, rotation=45)
plt.xlabel('Sensor', fontsize=12)
plt.ylabel('Variance', fontsize=12)
plt.title('Sensor Variance Analysis (Informative vs Constant)', fontsize=14,
        fontweight='bold')
plt.axhline(y=0.01, color='red', linestyle='--', linewidth=2, label='Threshold (0.01)')
plt.legend()
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

print(f"\nSensors with variance > 0.01: {(sensor_variances > 0.01).sum()}")
print(f"Sensors with variance <= 0.01: {(sensor_variances <= 0.01).sum()}")
```



Sensors with variance > 0.01: 11
 Sensors with variance <= 0.01: 10

1.7 Summary

Key Insights:

- 100 training engines with varying lifecycles (128-362 cycles)
- RUL capped at 125 cycles (standard practice in prognostics)
- 10 sensors removed — near-zero variance, constant on FD001
- 11 informative sensors retained for feature engineering
- Rolling window features capture temporal degradation patterns
- Data pipeline validated and ready for XGBoost training